

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.1      v tibble    3.2.1
## v lubridate  1.9.4      v tidyr     1.3.1
## v purrr      1.0.4
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(PlayerRatings)
library(BradleyTerry2)
library(dplyr)
library(tidyr)
library(ggplot2)
```

```
matches <- read.csv("atp_matches_with_features.csv")

# remove matches with carpet surface
matches <- matches %>% filter(surface != "Carpet")

# sort by date
matches$tourney_date <- as.Date(matches$tourney_date)
matches <- matches[order(matches$tourney_date), ]

# identify all winner rolling columns
w_rolling_cols <- grep("^w_rolling_10_", colnames(matches), value = TRUE)

# create difference columns dynamically
for (col in w_rolling_cols) {
  base_stat <- sub("^w_rolling_10_", "", col)
  l_col <- paste0("l_rolling_10_", base_stat)

  if (l_col %in% colnames(matches)) {
    diff_col <- paste0("diff_rolling_10_", base_stat)
    matches[[diff_col]] <- matches[[col]] - matches[[l_col]]
  }
}

# diff feature
matches$diff_h2h = matches$w_previous_wins - matches$l_previous_wins

# win loss for BT
matches$win <- 1L
matches$loss <- 0L
```

```
### ADD ELO ###
```

```
# 0. Make sure your full `matches` is sorted by date
```

```

matches <- matches[order(matches$tournament_date), ]

# 1. Initialize
initial_rating <- 1500
k_factor <- 20
all_players <- unique(c(matches$winner_name, matches$loser_name))
# start everybody at 1500
current_ratings <- setNames(
  rep(initial_rating, length(all_players)),
  all_players
)

# 2. Pre-allocate three new columns
matches$elo_w <- NA_real_ # winner's rating *before* the match
matches$elo_l <- NA_real_ # loser's rating *before* the match
matches$elo_prob <- NA_real_ # predicted win prob from Elo
matches$diff_elo <- NA_real_ # to enrich BT

# 3. helper for win-prob
calculate_win_prob <- function(diff) 1 / (1 + 10^(-diff/400))

# 4. loop through every match in chronological order
for (i in seq_len(nrow(matches))) {
  w <- matches$winner_name[i]
  l <- matches$loser_name[i]

  # 4a. get pre-match ratings (default 1500 if brand-new)
  r_w <- current_ratings[w]
  r_l <- current_ratings[l]

  # 4b. record them
  matches$elo_w[i] <- r_w
  matches$elo_l[i] <- r_l

  # 4c. compute predicted win-prob
  p_win <- calculate_win_prob(r_w - r_l)
  matches$elo_prob[i] <- p_win
  matches$diff_elo[i] <- r_w - r_l

  # 4d. update with outcome
  current_ratings[w] <- r_w + k_factor * (1 - p_win)
  current_ratings[l] <- r_l + k_factor * (0 - (1 - p_win))
}

### TRAN TEST SPLIT ###
matches$season <- as.numeric(format(matches$tournament_date, "%Y"))
train_matches <- matches %>% filter(season < 2024) %>% filter(season >= 2019)
test_matches <- matches %>% filter(season == 2024)
cat("Training set matches:", nrow(train_matches), "\n")

```

```
## Training set matches: 11606
```

```
cat("Testing set matches:", nrow(test_matches), "\n")
```

```
## Testing set matches: 2839
```

```
# filter and drop anyone who has only won or only lost in dataset
```

```
filter_pure_players <- function(df) {
  wins <- df %>%
    count(player = winner_name, name = "n_wins")

  losses <- df %>%
    count(player = loser_name, name = "n_losses")

  counts <- full_join(wins, losses, by = "player") %>%
    replace_na(list(n_wins = 0, n_losses = 0)) %>%
    mutate(n_total = n_wins + n_losses)

  bad_players <- counts %>%
    filter(n_wins == 0 | n_losses == 0) %>%
    filter(n_total <= 10) %>%
    pull(player)

  df %>%
    filter(
      !as.character(winner_name) %in% bad_players,
      !as.character(loser_name) %in% bad_players
    )
}

train_matches <- filter_pure_players(train_matches)
test_matches <- filter_pure_players(test_matches)
```

```
### ELO ###
```

```
# elo
```

```
train_games <- data.frame(
  Date = as.numeric(format(train_matches$tourney_date, "%Y%m%d")),
  Team1 = train_matches$winner_name,
  Team2 = train_matches$loser_name,
  Score = 1,
  stringsAsFactors = FALSE
)
elo_model_train <- elo(train_games, init = 1500, kfac = 20, gamma = 0)
```

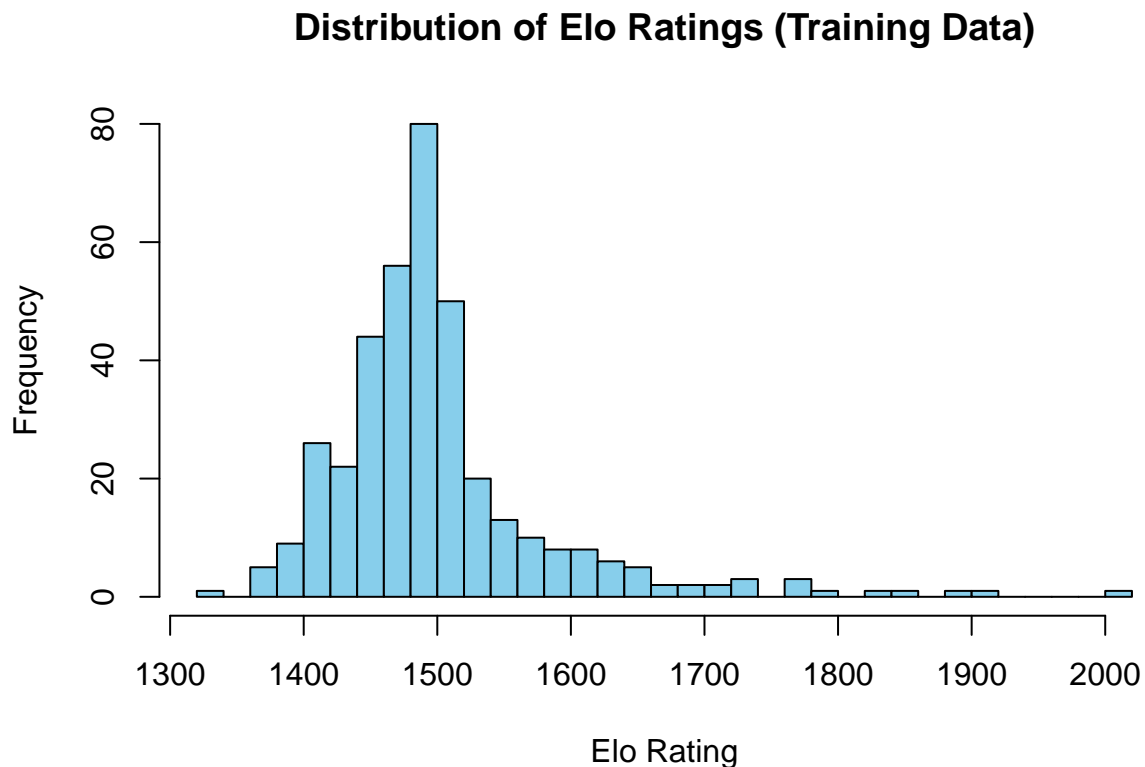
```
# top 10 players
```

```
elo_ratings_train <- elo_model_train$ratings
elo_ratings_train <- elo_ratings_train[order(elo_ratings_train$Rating, decreasing = TRUE), ]
top10_elo_train <- head(elo_ratings_train, 10)
print(top10_elo_train)
```

```
##           Player   Rating Games Win Draw Loss Lag
## 1   Novak Djokovic 2017.002  267 234    0  33   2
## 2    Jannik Sinner 1916.464  255 185    0  70   1
```

```
## 3    Carlos Alcaraz 1894.688   192 150    0   42    7
## 4    Daniil Medvedev 1849.723   326 247    0   79    7
## 5      Rafael Nadal 1820.167   175 145    0   30   47
## 6    Alexander Zverev 1795.039   289 203    0   86    7
## 7      Andrey Rublev 1769.745   327 234    0   93    7
## 8    Grigor Dimitrov 1765.806   216 124    0   92    9
## 9      Nick Kyrgios 1760.050   101  69    0   32   32
## 10   Hubert Hurkacz 1736.541   258 157    0  101    7
```

```
# plot distribution of Elo ratings based on training data
hist(elo_ratings_train$Rating,
     main = "Distribution of Elo Ratings (Training Data)",
     xlab = "Elo Rating",
     col = "skyblue",
     breaks = 30)
```



```
# evaluation metrics
correct_prediction <- test_matches$elo_prob > 0.5
accuracy <- mean(correct_prediction, na.rm = TRUE)
misclassification_rate <- 1 - accuracy
cat("Misclassification Rate:", misclassification_rate, "\n")
```

```
## Misclassification Rate: 0.3690078
```

```

brier_score <- (1 - test_matches$elo_prob)^2
mean_brier_score <- mean(brier_score, na.rm = TRUE)
cat("Mean Brier Score:", mean_brier_score, "\n")

```

```
## Mean Brier Score: 0.2204229
```

```

# look at elo's biggest fails
worst_elo <- matches %>%
  # only consider rows where we actually computed elo_prob
  filter(!is.na(elo_prob)) %>%
  # sort by ascending elo_prob
  arrange(elo_prob) %>%
  # take the 10 lowest
  slice_head(n = 10)

print(worst_elo)

```

```

##   tourney_id      tourney_name surface draw_size tourney_level tourney_date
## 1  2016-0410  Monte Carlo Masters    Clay         64             M  2016-04-11
## 2  2017-580   Australian Open      Hard        128             G  2017-01-16
## 3  2014-500                Halle    Grass         28             A  2014-06-09
## 4  2018-M006 Indian Wells Masters    Hard        128             M  2018-03-05
## 5  2014-540                Wimbledon Grass        128             G  2014-06-23
## 6  2017-0495                Dubai    Hard         32             A  2017-02-27
## 7  2024-0404 Indian Wells Masters    Hard        128             M  2024-03-04
## 8  2016-540                Wimbledon Grass        128             G  2016-06-27
## 9  2016-0495                Dubai    Hard         32             A  2016-02-22
## 10 2013-505                Santiago Clay         28             A  2013-02-04
##   match_num winner_id winner_seed winner_entry winner_name winner_hand
## 1         285   106210           NA           Jiri Vesely             L
## 2         195   104797           NA           Denis Istomin            R
## 3          13   104460           NA            WC Dustin Brown            R
## 4         242   106121           NA            Q  Taro Daniel            R
## 5         120   106401           NA            WC  Nick Kyrgios            R
## 6         291   105539           NA            Q  Evgeny Donskoy            R
## 7         285   208134           NA            LL   Luca Nardi            R
## 8         196   105023           28           Sam Querrey            R
## 9         297   103852            6           Feliciano Lopez            L
## 10        27   104547           NA           Horacio Zeballos            L
##   winner_ioc loser_id loser_seed loser_entry loser_name loser_hand
## 1          CZE  104925            1      Novak Djokovic            R
## 2          UZB  104925            2      Novak Djokovic            R
## 3          JAM  104745            1      Rafael Nadal             L
## 4          JPN  104925           10      Novak Djokovic            R
## 5          AUS  104745            2      Rafael Nadal             L
## 6          RUS  103819            3      Roger Federer            R
## 7          ITA  104925            1      Novak Djokovic            R
## 8          USA  104925            1      Novak Djokovic            R
## 9          ESP  104925            1      Novak Djokovic            R
## 10         ARG  104745            1      Rafael Nadal             L
##   loser_ioc      score best_of round minutes w_ace w_df w_svpt
## 1         SRB      6-4 2-6 6-4          3  R32    126      4      4    76
## 2         SRB 7-6(8) 5-7 2-6 7-6(5) 6-4          5  R64    288     17      3   194

```

## 3	ESP		6-4	6-1	3	R16	60	11	5	51	
## 4	SRB		7-6(3)	4-6	6-1	3	R64	150	1	8	107
## 5	ESP	7-6(5)	5-7	7-6(5)	6-3	5	R16	178	37	4	144
## 6	SUI	3-6	7-6(7)	7-6(5)		3	R16	122	6	3	95
## 7	SRB		6-4	3-6	6-3	3	R32	140	6	7	90
## 8	SRB	7-6(6)	6-1	3-6	7-6(5)	5	R32	176	31	2	160
## 9	SRB		6-3	0-0	RET	3	QF	33	2	2	22
## 10	ESP	6-7(2)	7-6(6)	6-4		3	F	167	12	0	96
##	w_1stIn	w_1stWon	w_2ndWon	w_SvGms	w_bpSaved	w_bpFaced	l_ace	l_df	l_svpt		
## 1	45	31	19	14	1	4	3	3	92		
## 2	134	94	27	27	9	15	14	9	185		
## 3	31	26	12	9	1	1	2	0	47		
## 4	52	37	28	15	5	7	4	4	94		
## 5	98	81	25	23	2	3	11	3	137		
## 6	46	35	29	16	1	5	12	2	109		
## 7	64	44	13	14	2	4	4	1	95		
## 8	100	79	25	20	14	17	7	2	146		
## 9	10	9	8	4	0	0	1	1	29		
## 10	58	45	27	17	2	3	4	1	106		
##	l_1stIn	l_1stWon	l_2ndWon	l_SvGms	l_bpSaved	l_bpFaced	winner_rank	points			
## 1	61	44	14	14	5	8		840			
## 2	107	77	43	27	10	14		498			
## 3	32	22	6	8	6	9		633			
## 4	60	39	18	14	8	11		539			
## 5	100	77	22	22	3	4		400			
## 6	80	59	15	17	1	4		493			
## 7	57	44	16	14	8	11		513			
## 8	97	63	27	20	11	15		1075			
## 9	16	11	6	5	4	6		1495			
## 10	77	57	18	17	5	7		708			
##	loser_rank	points	w_rank	l_rank	w_ht	l_ht	w_age	l_age			
## 1		16540	55	1	198	188	22.7	28.8			
## 2		11780	117	2	185	188	30.3	29.6			
## 3		12500	85	1	196	185	29.5	28.0			
## 4		2380	109	13	191	188	25.1	30.7			
## 5		12500	144	1	193	185	19.1	28.0			
## 6		3260	116	10	185	185	26.8	35.5			
## 7		9675	123	1	185	188	20.5	36.7			
## 8		16950	41	1	198	188	28.7	29.1			
## 9		16790	24	1	188	188	34.4	28.7			
## 10		5400	73	5	188	185	27.7	26.6			
##	w_rolling_10_ace	per_svgm	w_rolling_10_df	per_svgm							
## 1		0.4511491		0.2078711							
## 2		0.5833922		0.1660722							
## 3		0.6999999		0.3799950							
## 4		0.2976221		0.4324123							
## 5		0.8449815		0.2864957							
## 6		0.3989322		0.1825722							
## 7		0.2319444		0.1845635							
## 8		1.1158604		0.1560408							
## 9		1.0280240		0.3345285							
## 10		0.4224870		0.1723922							
##	w_rolling_10_svpt	per_svgm	w_rolling_10_1stIn	per_svgm							
## 1		6.660308		4.318270							

## 2	6.863417	4.349540	
## 3	6.493274	4.030144	
## 4	6.852013	3.966637	
## 5	6.696712	4.315922	
## 6	6.757023	4.066698	
## 7	6.253968	4.262897	
## 8	5.920988	3.569351	
## 9	6.394547	3.462069	
## 10	6.441168	3.999141	
##	w_rolling_10_1stWon_per_svgm	w_rolling_10_2ndWon_per_svgm	w_rolling_10_SvGms
## 1	2.982642	1.1922912	12.4
## 2	2.997867	1.0893146	15.7
## 3	3.065017	1.1948563	12.0
## 4	2.579143	1.2825278	14.9
## 5	3.153023	1.1751100	20.0
## 6	2.802148	1.3225900	14.0
## 7	2.752540	0.8684921	11.7
## 8	3.053470	1.2897872	14.5
## 9	2.754359	1.5197981	15.1
## 10	2.750759	1.2496695	12.3
##	w_rolling_10_bpSaved_per_svgm	w_rolling_10_bpSaved_per_bpFaced_per_svgm	
## 1	0.4206411	0.05825708	
## 2	0.4697678	0.06386140	
## 3	0.2225720	0.05620622	
## 4	0.4987574	0.04407812	
## 5	0.3130449	0.03015992	
## 6	0.4700900	0.04824335	
## 7	0.2821429	0.04087693	
## 8	0.1862734	0.04731727	
## 9	0.2373069	0.04352061	
## 10	0.3053836	0.04631534	
##	w_rolling_10_bpFaced_per_svgm	w_rolling_10_rank	w_rolling_10_ht
## 1	0.6420977	54.7	198
## 2	0.7603549	107.8	185
## 3	0.4000875	90.1	196
## 4	0.8251344	108.2	191
## 5	0.4938304	174.7	193
## 6	0.7230115	107.6	185
## 7	0.6248413	135.4	185
## 8	0.2471916	37.3	198
## 9	0.3702483	21.5	188
## 10	0.5366593	90.6	188
##	w_rolling_10_age	l_rolling_10_ace_per_svgm	l_rolling_10_df_per_svgm
## 1	22.65	0.34693001	0.26968975
## 2	30.01	0.24847375	0.18970238
## 3	29.37	0.08620614	0.08698878
## 4	24.86	0.38599512	0.23354701
## 5	18.78	0.21323880	0.06707251
## 6	26.54	0.76088593	0.12058827
## 7	20.05	0.63787879	0.13021645
## 8	28.66	0.22706855	0.16498849
## 9	34.33	0.49201034	0.13419279
## 10	27.46	0.25666616	0.12407466
##	l_rolling_10_svpt_per_svgm	l_rolling_10_1stIn_per_svgm	

## 1	6.491068	4.234091	
## 2	6.376813	4.222031	
## 3	6.050694	4.278675	
## 4	6.153907	4.198639	
## 5	5.729080	3.930226	
## 6	5.914489	3.781874	
## 7	6.087828	3.863506	
## 8	6.297842	4.318990	
## 9	5.879572	3.756680	
## 10	6.097670	4.094065	
##	l_rolling_10_1stWon_per_svgm	l_rolling_10_2ndWon_per_svgm	l_rolling_10_SvGms
## 1	3.083806	1.259513	9.6
## 2	3.162063	1.151723	11.6
## 3	3.006298	1.066767	13.3
## 4	3.150263	1.076288	12.7
## 5	2.975039	1.111718	14.4
## 6	2.988445	1.193301	19.4
## 7	2.963752	1.232626	15.1
## 8	3.015420	1.138755	15.6
## 9	2.853129	1.325123	13.8
## 10	3.086653	1.210515	12.8
##	l_rolling_10_bpSaved_per_svgm	l_rolling_10_bpSaved_per_bpFaced_per_svgm	
## 1	0.2527525	0.04460017	
## 2	0.2712042	0.06716106	
## 3	0.3279863	0.06085844	
## 4	0.1688462	0.04599953	
## 5	0.2789513	0.05712679	
## 6	0.1885218	0.02854249	
## 7	0.2391775	0.04660534	
## 8	0.3065307	0.03873206	
## 9	0.2631004	0.06151036	
## 10	0.3097921	0.06417177	
##	l_rolling_10_bpFaced_per_svgm	l_rolling_10_rank	l_rolling_10_ht
## 1	0.3746212	1.0	188
## 2	0.3991010	2.0	188
## 3	0.4691441	1.0	185
## 4	0.2641575	8.0	188
## 5	0.3880318	1.0	185
## 6	0.2938705	13.5	185
## 7	0.3626696	1.0	188
## 8	0.4592503	1.0	188
## 9	0.3654388	1.0	188
## 10	0.3981124	2.9	185
##	l_rolling_10_age	w_previous_wins	l_previous_wins
## 1	28.76	0	0
## 2	29.52	0	5
## 3	27.90	0	0
## 4	30.29	0	0
## 5	27.94	0	0
## 6	35.29	0	0
## 7	36.61	0	0
## 8	29.01	1	6
## 9	28.62	0	7
## 10	26.15	0	1


```

##      diff_rolling_10_ace_per_svgm diff_rolling_10_df_per_svgm
## 1              0.1042191          -0.061818606
## 2              0.3349185          -0.023630151
## 3              0.6137938           0.293006248
## 4             -0.0883730           0.198865341
## 5              0.6317427           0.219423221
## 6             -0.3619537           0.061983885
## 7             -0.4059343           0.054347042
## 8              0.8887918          -0.008947725
## 9              0.5360136           0.200335699
## 10             0.1658209           0.048317531
##      diff_rolling_10_svpt_per_svgm diff_rolling_10_1stIn_per_svgm
## 1              0.1692403           0.08417876
## 2              0.4866037           0.12750809
## 3              0.4425795          -0.24853144
## 4              0.6981059          -0.23200129
## 5              0.9676317           0.38569576
## 6              0.8425337           0.28482441
## 7              0.1661400           0.39939033
## 8             -0.3768531          -0.74963969
## 9              0.5149756          -0.29461065
## 10             0.3434972          -0.09492438
##      diff_rolling_10_1stWon_per_svgm diff_rolling_10_2ndWon_per_svgm
## 1             -0.10116395          -0.06722180
## 2             -0.16419622          -0.06240850
## 3              0.05871903           0.12808939
## 4             -0.57111932           0.20623962
## 5              0.17798370           0.06339159
## 6             -0.18629632           0.12928936
## 7             -0.21121212          -0.36413420
## 8              0.03805058           0.15103180
## 9             -0.09876932           0.19467527
## 10            -0.33589437           0.03915466
##      diff_rolling_10_SvGms diff_rolling_10_bpSaved_per_svgm
## 1              2.8              0.16788855
## 2              4.1              0.19856364
## 3             -1.3             -0.10541433
## 4              2.2              0.32991124
## 5              5.6              0.03409359
## 6             -5.4              0.28156816
## 7             -3.4              0.04296537
## 8             -1.1             -0.12025727
## 9              1.3             -0.02579346
## 10            -0.5             -0.00440852
##      diff_rolling_10_bpSaved_per_bpFaced_per_svgm
## 1              0.013656907
## 2             -0.003299658
## 3             -0.004652223
## 4             -0.001921409
## 5             -0.026966869
## 6              0.019700855
## 7             -0.005728407
## 8              0.008585212
## 9             -0.017989751

```

```
## 10                                -0.017856426
##   diff_rolling_10_bpFaced_per_svgm diff_rolling_10_rank diff_rolling_10_ht
## 1                                0.26747645                53.7                10
## 2                                0.36125381                105.8               -3
## 3                               -0.06905666                 89.1                11
## 4                                0.56097688                100.2                 3
## 5                                0.10579867                173.7                 8
## 6                                0.42914101                 94.1                 0
## 7                                0.26217172                134.4               -3
## 8                               -0.21205876                 36.3                10
## 9                                0.00480957                 20.5                 0
## 10                               0.13854693                 87.7                 3
##   diff_rolling_10_age diff_h2h win loss   elo_w   elo_l   elo_prob   diff_elo
## 1                -6.11      0   1   0 1541.692 2289.451 0.01332828 -747.7595
## 2                 0.49     -5   1   0 1551.530 2221.740 0.02067291 -670.2104
## 3                 1.47      0   1   0 1527.063 2181.300 0.02261885 -654.2374
## 4                -5.43      0   1   0 1459.490 2108.049 0.02335290 -648.5587
## 5                -9.16      0   1   0 1515.022 2162.449 0.02350193 -647.4271
## 6                -8.75      0   1   0 1449.363 2095.363 0.02369118 -646.0002
## 7               -16.56      0   1   0 1460.460 2102.475 0.02422768 -642.0146
## 8                -0.35     -5   1   0 1664.010 2284.417 0.02734879 -620.4076
## 9                 5.71     -7   1   0 1684.826 2291.480 0.02953551 -606.6541
## 10                1.31     -1   1   0 1491.265 2097.051 0.02967899 -605.7865
##   season
## 1    2016
## 2    2017
## 3    2014
## 4    2018
## 5    2014
## 6    2017
## 7    2024
## 8    2016
## 9    2016
## 10   2013
```

```
extract_diff_coefficients <- function(model, pattern = "^diff", sort_by = "p.value", sort = TRUE) {
  coef_table <- summary(model)$coefficients
  matching_rows <- grep(pattern, rownames(coef_table))
  results <- data.frame(
    term      = rownames(coef_table)[matching_rows],
    estimate  = coef_table[matching_rows, "Estimate"],
    p.value   = coef_table[matching_rows, "Pr(>|z|)"],
    row.names = NULL
  )
  if (sort) {
    results <- results[order(results[[sort_by]], decreasing = FALSE), ]
  }
  results
}

evaluate_glm_model <- function(model, test_data, newdata, valid_idx, prediction_col = "pred_prob", thresh = 0.05) {
  preds <- predict(
    model,
    newdata = newdata,
```

```

    type = "response"
  )

  test_data[[prediction_col]] <- NA_real_
  test_data[[prediction_col]][valid_idx] <- preds

  correct_prediction <- test_data[[prediction_col]] > threshold
  accuracy <- mean(correct_prediction, na.rm = TRUE)
  misclassification_rate <- 1 - accuracy

  brier_score <- (1 - test_data[[prediction_col]])^2
  mean_brier_score <- mean(brier_score, na.rm = TRUE)

  cat("Evaluation Results:\n")
  cat("-----\n")
  cat("Misclassification Rate:", round(misclassification_rate, 4), "\n")
  cat("Mean Brier Score:", round(mean_brier_score, 4), "\n")

  # Return results as a list
  results <- list(
    predictions = test_data[[prediction_col]],
    accuracy = accuracy,
    misclassification_rate = misclassification_rate,
    mean_brier_score = mean_brier_score,
    test_data = test_data
  )

  return(results)
}

# train levels
all_players <- (unique(c(train_matches$winner_name,
                        train_matches$loser_name)))

train_matches$winner_name = factor(train_matches$winner_name, levels = all_players)
train_matches$loser_name = factor(train_matches$loser_name, levels = all_players)

# re-level the test set to exactly those levels used in train
test_matches$winner_name = factor(test_matches$winner_name, levels = all_players)
test_matches$loser_name = factor(test_matches$loser_name, levels = all_players)

# drop any rows that now have NA (i.e. brand-new players unseen in train)
valid_idx <- which(
  !is.na(test_matches$winner_name) &
  !is.na(test_matches$loser_name)
)

test_valid <- test_matches[valid_idx, ]
nrow(test_valid)

```

```
## [1] 2691
```

```
nrow(test_valid)
```

```
## [1] 2459
```

```
#### VANILLA BT ####
```

```
X <- model.matrix(~ train_matches$winner_name - 1) - model.matrix(~ train_matches$loser_name - 1)
```

```
model <- glm(train_matches$win ~ X[, -1] - 1, family = binomial, data = train_matches)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
ratings <- coef(model)
```

```
ratings_df <- data.frame(  
  player = names(ratings),  
  rating = ratings  
)
```

```
ratings_df$player <- sub("^.*winner_name", "", ratings_df$player)
```

```
sorted_ratings_df <- ratings_df[order(ratings_df$rating, decreasing = TRUE), ]  
top10 <- sorted_ratings_df[1:10, ]
```

```
# Print the top 10 rated players.
```

```
cat("Top 10 Rated Players (Higher rating = stronger):\n")
```

```
## Top 10 Rated Players (Higher rating = stronger):
```

```
print(top10)
```

```
##
```

	player	rating
## X[, -1]train_matches\$winner_name	Nick Hardt	13.913203
## X[, -1]train_matches\$winner_name	Novak Djokovic	2.656996
## X[, -1]train_matches\$winner_name	Rafael Nadal	2.248558
## X[, -1]train_matches\$winner_name	Roger Federer	1.950682
## X[, -1]train_matches\$winner_name	Carlos Alcaraz	1.800364
## X[, -1]train_matches\$winner_name	Daniil Medvedev	1.720706
## X[, -1]train_matches\$winner_name	Franco Agamenone	1.445257
## X[, -1]train_matches\$winner_name	Jannik Sinner	1.426718
## X[, -1]train_matches\$winner_name	Stefanos Tsitsipas	1.401950
## X[, -1]train_matches\$winner_name	Alexander Zverev	1.377390

```
# test
```

```
predict_probs <- function(tourney_df, ratings_df) {  
  tourney_df %>%  
    left_join(ratings_df, by = c("winner_name" = "player")) %>%  
    rename(W_rating = rating) %>%  
    left_join(ratings_df, by = c("loser_name" = "player")) %>%  
    rename(L_rating = rating) %>%  
    mutate(  
      pred_prob = exp(W_rating) / (exp(W_rating) + exp(L_rating))  
    )  
}
```

```

}

pred_prob <- predict_probs(test_matches, ratings_df)$pred_prob

# evaluation metrics
correct_prediction <- pred_prob > 0.5
accuracy <- mean(correct_prediction, na.rm = TRUE)
misclassification_rate <- 1 - accuracy
cat("Misclassification Rate:", misclassification_rate, "\n")

```

```
## Misclassification Rate: 0.3926658
```

```

brier_score <- (1 - pred_prob)^2
mean_brier_score <- mean(brier_score, na.rm = TRUE)
cat("Mean Brier Score:", mean_brier_score, "\n")

```

```
## Mean Brier Score: 0.231361
```

```

### BT WITH COVARIATES ###
covariates <- c(
  "diff_rolling_10_ace_per_svgm",
  "diff_rolling_10_df_per_svgm",
  "diff_rolling_10_svpt_per_svgm",
  "diff_rolling_10_1stIn_per_svgm",
  "diff_rolling_10_1stWon_per_svgm",
  "diff_rolling_10_2ndWon_per_svgm",
  "diff_rolling_10_SvGms",
  "diff_rolling_10_bpSaved_per_svgm",
  "diff_rolling_10_bpSaved_per_bpFaced_per_svgm",
  "diff_rolling_10_bpFaced_per_svgm",
  "diff_rolling_10_rank",
  "diff_rolling_10_ht",
  "diff_rolling_10_age",
  "diff_h2h"
)
cov_str <- paste(covariates, collapse = " + ")

model_w_diffs <- glm(
  as.formula(
    paste0(
      "win ~ X[, -1] + ",
      cov_str,
      " - 1"
    )
  ),
  data = train_matches,
  family = binomial(link = "logit")
)

```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
diff_coefs <- extract_diff_coefficients(model_w_diffs)
diff_coefs
```

```
##              term      estimate    p.value
## 12      diff_rolling_10_age  1.343533097 6.529092e-51
## 11      diff_rolling_10_rank  0.001749588 1.910604e-06
## 2       diff_rolling_10_df_per_svgm -1.322543348 5.586404e-06
## 5       diff_rolling_10_1stWon_per_svgm  0.965788475 9.297474e-03
## 9  diff_rolling_10_bpSaved_per_bpFaced_per_svgm  4.908493101 5.818879e-02
## 10      diff_rolling_10_bpFaced_per_svgm -1.561801088 1.333007e-01
## 6       diff_rolling_10_2ndWon_per_svgm  0.576142635 1.446710e-01
## 3       diff_rolling_10_svpt_per_svgm -0.275939123 1.870595e-01
## 4       diff_rolling_10_1stIn_per_svgm -0.237914306 2.105609e-01
## 13              diff_h2h    0.018013770 2.894171e-01
## 8       diff_rolling_10_bpSaved_per_svgm  0.993491910 3.878285e-01
## 7       diff_rolling_10_SvGms  0.010429035 4.835506e-01
## 1       diff_rolling_10_ace_per_svgm -0.064999450 6.985835e-01
```

```
# combine into one newdata frame
X <- model.matrix(~ test_valid$winner_name - 1) - model.matrix(~ test_valid$loser_name - 1)
newdata <- cbind(
  as.data.frame(X),
  test_valid[, covariates, drop = FALSE]
)

results <- evaluate_glm_model(
  model = model_w_diffs,
  test_data = test_matches,
  newdata = newdata,
  valid_idx = valid_idx
)
```

```
## Evaluation Results:
## -----
## Misclassification Rate: 0.342
## Mean Brier Score: 0.2168
```

```
### BT WITH COVARIATES REDUCED ###
covariates <- c(
  "diff_rolling_10_ace_per_svgm",
  "diff_rolling_10_df_per_svgm",
  "diff_rolling_10_svpt_per_svgm",
  "diff_rolling_10_1stIn_per_svgm",
  "diff_rolling_10_1stWon_per_svgm",
  "diff_rolling_10_2ndWon_per_svgm",
  "diff_rolling_10_SvGms",
  "diff_rolling_10_bpSaved_per_svgm",
  "diff_rolling_10_bpSaved_per_bpFaced_per_svgm",
  "diff_rolling_10_bpFaced_per_svgm",
  "diff_rolling_10_rank",
  "diff_rolling_10_ht",
  "diff_rolling_10_age",

```

```

"diff_h2h"
)
cov_str <- paste(covariates, collapse = " + ")

X <- model.matrix(~ train_matches$winner_name - 1) - model.matrix(~ train_matches$loser_name - 1)
model_w_diffs_red <- glm(
  as.formula(
    paste0(
      "win ~ X[, -1] + ",
      cov_str,
      " - 1"
    )
  ),
  data = train_matches,
  family = binomial(link = "logit")
)

```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```

diff_coefs <- extract_diff_coefficients(model_w_diffs_red, sort=TRUE)
diff_coefs

```

```

##              term      estimate      p.value
## 7              diff_rolling_10_age  1.340788829 7.240926e-51
## 3              diff_rolling_10_1stWon_per_svgm  1.458405970 2.889350e-31
## 2              diff_rolling_10_svpt_per_svgm -0.780156579 1.533684e-25
## 4              diff_rolling_10_2ndWon_per_svgm  1.396246964 1.353100e-18
## 6              diff_rolling_10_rank   0.001762834 1.585258e-06
## 1              diff_rolling_10_df_per_svgm -1.170818524 1.140837e-05
## 5 diff_rolling_10_bpSaved_per_bpFaced_per_svgm  2.777288267 1.109569e-01
## 8              diff_h2h   0.018105993 2.865224e-01

```

```

# combine into one newdata frame
X <- model.matrix(~ test_valid$winner_name - 1) - model.matrix(~ test_valid$loser_name - 1)
newdata <- cbind(
  as.data.frame(X),
  test_valid[, covariates, drop = FALSE]
)

results <- evaluate_glm_model(
  model = model_w_diffs_red,
  test_data = test_matches,
  newdata = newdata,
  valid_idx = valid_idx
)

```

```

## Evaluation Results:
## -----
## Misclassification Rate: 0.34
## Mean Brier Score: 0.217

```

```

### BT WITH COVARIATES REDUCED BY SURFACE ###
X <- model.matrix(~ train_matches$winner_name - 1) - model.matrix(~ train_matches$loser_name - 1)
model_w_diffs_surface <- glm(
  as.formula(
    paste0(
      "win ~ X[, -1] + (",
      cov_str,
      ") : surface - 1"
    )
  ),
  data = train_matches,
  family = binomial(link = "logit")
)

```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```

diff_coefs <- extract_diff_coefficients(model_w_diffs_surface, sort=FALSE)
diff_coefs

```

	term	estimate
## 1	diff_rolling_10_df_per_svgm:surfaceClay	-1.877282046
## 2	diff_rolling_10_df_per_svgm:surfaceGrass	-0.143376301
## 3	diff_rolling_10_df_per_svgm:surfaceHard	-1.031402739
## 4	diff_rolling_10_svpt_per_svgm:surfaceClay	-0.392713263
## 5	diff_rolling_10_svpt_per_svgm:surfaceGrass	-1.728070922
## 6	diff_rolling_10_svpt_per_svgm:surfaceHard	-0.809332382
## 7	diff_rolling_10_1stWon_per_svgm:surfaceClay	0.681850898
## 8	diff_rolling_10_1stWon_per_svgm:surfaceGrass	3.176587214
## 9	diff_rolling_10_1stWon_per_svgm:surfaceHard	1.527768419
## 10	diff_rolling_10_2ndWon_per_svgm:surfaceClay	0.599140796
## 11	diff_rolling_10_2ndWon_per_svgm:surfaceGrass	3.137077719
## 12	diff_rolling_10_2ndWon_per_svgm:surfaceHard	1.521025162
## 13	diff_rolling_10_bpSaved_per_bpFaced_per_svgm:surfaceClay	1.030871487
## 14	diff_rolling_10_bpSaved_per_bpFaced_per_svgm:surfaceGrass	-8.214851229
## 15	diff_rolling_10_bpSaved_per_bpFaced_per_svgm:surfaceHard	5.465138868
## 16	diff_rolling_10_rank:surfaceClay	0.001252076
## 17	diff_rolling_10_rank:surfaceGrass	0.002943580
## 18	diff_rolling_10_rank:surfaceHard	0.001658732
## 19	diff_rolling_10_ht:surfaceClay	0.003018123
## 20	diff_rolling_10_ht:surfaceGrass	-0.011178967
## 21	diff_rolling_10_age:surfaceClay	1.341453092
## 22	diff_rolling_10_age:surfaceGrass	1.374128007
## 23	diff_rolling_10_age:surfaceHard	1.347935711
## 24	diff_h2h:surfaceClay	0.032464789
## 25	diff_h2h:surfaceGrass	-0.010039361
## 26	diff_h2h:surfaceHard	0.016768835
##	p.value	
## 1	7.522918e-07	
## 2	7.855126e-01	
## 3	6.161771e-04	
## 4	9.507736e-04	
## 5	2.530505e-15	
## 6	2.808953e-18	


```
## 7 9.027279e-04
## 8 1.844799e-19
## 9 6.307114e-23
## 10 1.221952e-02
## 11 6.063637e-14
## 12 4.767100e-15
## 13 7.385806e-01
## 14 1.322305e-01
## 15 1.229408e-02
## 16 1.739045e-02
## 17 6.235364e-04
## 18 1.400067e-04
## 19 6.245040e-01
## 20 2.362331e-01
## 21 4.158398e-50
## 22 2.769495e-52
## 23 8.094602e-51
## 24 2.912763e-01
## 25 8.504903e-01
## 26 4.254454e-01
```

```
# combine into one newdata frame
X <- model.matrix(~ test_valid$winner_name - 1) - model.matrix(~ test_valid$loser_name - 1)
newdata <- cbind(
  as.data.frame(model.matrix(~ test_valid$winner_name - 1) - model.matrix(~ test_valid$loser_name - 1)),
  test_valid[, covariates, drop = FALSE],
  surface = test_valid$surface
)

results <- evaluate_glm_model(
  model = model_w_diffs_surface,
  test_data = test_matches,
  newdata = newdata,
  valid_idx = valid_idx
)
```

```
## Evaluation Results:
## -----
## Misclassification Rate: 0.3449
## Mean Brier Score: 0.2163
```

```
### EXAMINE STAT DIFFERENCES BETWEEN TOP GROUP VS WORST GROUP ###
# 1. pick the vars to test
test_vars <- covariates

# 2. label each match as 'Upset' or 'FavoriteWin'
cmp_df <- matches %>%
  filter(!is.na(elo_prob)) %>%
  mutate(
    result_group = case_when(
      elo_prob < 0.2 ~ "Upset",
      elo_prob > 0.8 ~ "FavoriteWin",
      TRUE ~ NA_character_
    )
  )
```

```

    )
  ) %>%
  filter(!is.na(result_group))

# 3. run t-tests for each variable
ttests <- map_dfr(test_vars, function(var){
  x <- pull(cmp_df, var)[cmp_df$result_group == "Upset"]
  y <- pull(cmp_df, var)[cmp_df$result_group == "FavoriteWin"]
  t <- t.test(y, x)
  tibble(
    stat = var,
    mean_fav = mean(y, na.rm=TRUE),
    mean_upset = mean(x, na.rm=TRUE),
    diff_means = mean(y, na.rm=TRUE) - mean(x, na.rm=TRUE),
    t_statistic = t$statistic,
    p_value = t$p.value
  )
}) %>%
  arrange(diff_means)

ttests <- ttests[order(ttests$p_value), ]
ttests

```

```
## # A tibble: 9 x 6
```

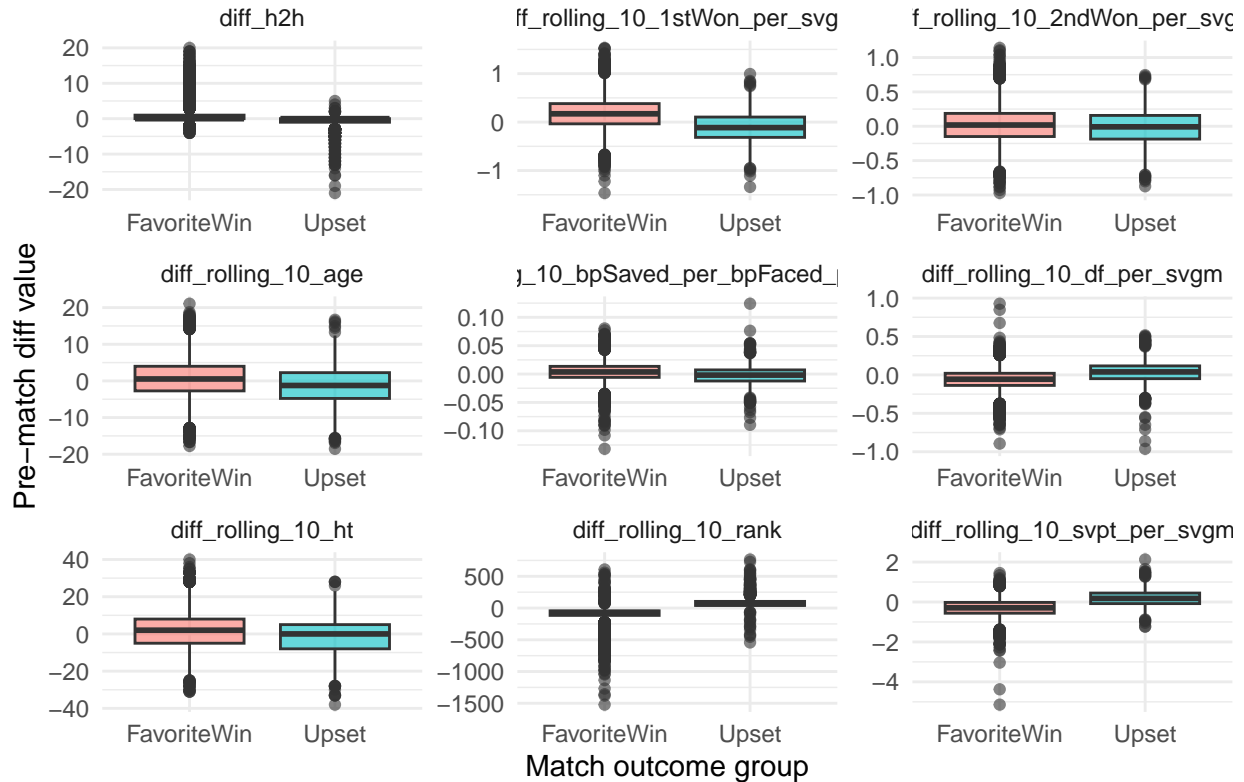
	stat	mean_fav	mean_upset	diff_means	t_statistic	p_value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	diff_rolling_10_rank	-9.80e+1	83.5	-181.	-80.9	0
## 2	diff_rolling_10_svpt_per~	-2.93e-1	0.188	-0.481	-46.2	0
## 3	diff_h2h	1.06e+0	-0.924	1.99	38.4	5.44e-256
## 4	diff_rolling_10_1stWon_p~	1.74e-1	-0.109	0.284	35.7	1.92e-226
## 5	diff_rolling_10_df_per_s~	-5.96e-2	0.0336	-0.0933	-26.6	2.47e-137
## 6	diff_rolling_10_bpSaved_~	3.94e-3	-0.00229	0.00624	15.3	1.63e- 50
## 7	diff_rolling_10_age	7.47e-1	-1.30	2.05	15.2	7.57e- 50
## 8	diff_rolling_10_ht	1.65e+0	-1.09	2.74	11.3	6.77e- 29
## 9	diff_rolling_10_2ndWon_p~	2.41e-2	-0.0155	0.0397	6.06	1.52e- 9

```

# 4. visualize distributions side-by-side
cmp_df %>%
  select(result_group, all_of(test_vars)) %>%
  pivot_longer(-result_group, names_to="stat", values_to="value") %>%
  ggplot(aes(x=result_group, y=value, fill=result_group)) +
  geom_boxplot(alpha=0.6) +
  facet_wrap(~stat, scales="free") +
  labs(
    x = "Match outcome group",
    y = "Pre-match diff value",
    title = "Upset vs Favorite-Win: Distributions of Key Diff Stats"
  ) +
  theme_minimal() +
  theme(legend.position="none")

```

Upset vs Favorite.Win: Distributions of Key Diff Stats



```
### DEPRECATED: TOO MUCH SPARSITY IN PLAYER-SURFACE BT ###
# 1. build the universe of all player*surface combos
all_players <- sort(unique(c(as.character(train_matches$winner_name),
                             as.character(train_matches$loser_name))))
surfaces <- sort(unique(as.character(train_matches$surface)))

all_surf_lvls <- as.vector(outer(all_players, surfaces,
                                FUN = function(p, s) paste(p, s, sep = ".")))

train_df <- train_matches %>%
  mutate(
    w_surf = factor(paste(winner_name, surface, sep = "."), levels = all_surf_lvls),
    l_surf = factor(paste(loser_name, surface, sep = "."), levels = all_surf_lvls)
  )

# 2. create winner*surface and loser*surface factors with those levels
train_df <- train_df %>%
  mutate(
    w_surf = factor(paste(winner_name, surface, sep = "."), levels = all_surf_lvls),
    l_surf = factor(paste(loser_name, surface, sep = "."), levels = all_surf_lvls)
  )

# 3. now build BT design matrices
Xw_surf <- model.matrix(~ w_surf - 1, data = train_df)
Xl_surf <- model.matrix(~ l_surf - 1, data = train_df)
```

```

X_bt_surf <- Xw_surf - Xl_surf # same columns in same order-you can subtract

# 4. bind with your diff features and fit
glm_input <- cbind(
  win = train_df$win,
  as.data.frame(X_bt_surf)
)

glm_surf <- glm(win ~ . - 1, data = glm_input, family = binomial)

## Warning: glm.fit: algorithm did not converge

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

# --- 1. Re-level test set to your training factors
test_matches <- test_matches %>%
  mutate(
    winner_name = factor(winner_name, levels = all_players),
    loser_name = factor(loser_name, levels = all_players),
    surface = factor(surface, levels = surfaces)
  )

# 2. Drop any matches with unseen players
valid_idx <- which(
  !is.na(test_matches$winner_name) &
  !is.na(test_matches$loser_name) &
  !is.na(test_matches$surface) # <- drop matches on unseen surfaces
)
test_valid <- test_matches[valid_idx, ]

# 3. Recreate the player*surface interaction on test
test_valid <- test_valid %>%
  mutate(
    w_surf = factor(paste(winner_name, surface, sep="."), levels = all_surf_lvls),
    l_surf = factor(paste(loser_name, surface, sep="."), levels = all_surf_lvls)
  )

# 4. Build the BT design matrix for player*surface
Xw_surf_test <- model.matrix(~ w_surf - 1, data = test_valid)
Xl_surf_test <- model.matrix(~ l_surf - 1, data = test_valid)
X_bt_surf_test <- Xw_surf_test - Xl_surf_test

# 5. Combine with your rolling-diff covariates
newdata <- cbind(
  as.data.frame(X_bt_surf_test)#,
  #test_valid[, covariates]
)

# 6. Predict win-probabilities
preds <- predict(
  glm_surf,
  newdata = newdata,

```

```

    type      = "response"
  )

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful cases

test_matches$pred_prob_bt_w_diffs_and_surf <- NA_real_
test_matches$pred_prob_bt_w_diffs_and_surf[valid_idx] <- preds

# evaluation metrics
test_matches$correct_prediction <- test_matches$pred_prob_bt_w_diffs_and_surf > 0.5
accuracy <- mean(test_matches$correct_prediction, na.rm = TRUE)
misclassification_rate <- 1 - accuracy
cat("Misclassification Rate:", misclassification_rate, "\n")

## Misclassification Rate: 0.4343229

test_matches$brier_score <- (1 - test_matches$pred_prob_bt_w_diffs_and_surf)^2
mean_brier_score <- mean(test_matches$brier_score, na.rm = TRUE)
cat("Mean Brier Score:", mean_brier_score, "\n")

## Mean Brier Score: 0.4337129

### RANDOM FOREST ###
library(randomForest)

## randomForest 4.7-1.2

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:dplyr':
##
##      combine

## The following object is masked from 'package:ggplot2':
##
##      margin

# 1. Identify your diff columns
diff_cols <- grep("^diff", names(train_matches), value = TRUE)

# 2. Build the RF training set
rf_pos <- train_matches %>%
  dplyr::select(all_of(diff_cols)) %>%
  mutate(label = 1)

```

```

rf_neg <- train_matches %>%
  dplyr::select(all_of(diff_cols)) %>%
  mutate(across(all_of(diff_cols), ~ -.), # flip the diffs
         label = 0)

rf_train <- bind_rows(rf_pos, rf_neg)

# 3. Fit a random forest
set.seed(42)
rf_mod <- randomForest(
  x      = rf_train %>% dplyr::select(all_of(diff_cols)),
  y      = factor(rf_train$label), # as a factor for classification
  ntree  = 500,
  importance = TRUE
)

print(rf_mod)

```

```

##
## Call:
## randomForest(x = rf_train %>% dplyr::select(all_of(diff_cols)), y = factor(rf_train$label), nt
##               Type of random forest: classification
##               Number of trees: 500
## No. of variables tried at each split: 3
##
## OOB estimate of error rate: 35.27%
## Confusion matrix:
##      0      1 class.error
## 0 7365 4034  0.3538907
## 1 4007 7392  0.3515221


```

```

varImpPlot(rf_mod) # see which diffs matter most

```

rf_mod



```
# 4. Predict on your test set
rf_test <- test_matches %>%
  filter(season > 2014) %>%
  dplyr::select(all_of(diff_cols))

# P(label=1) = probability that the "first-player" (i.e. winner_name) wins
test_matches$rf_prob <- predict(rf_mod, newdata = rf_test, type = "prob")[, "1"]

# 5. Evaluate
test_matches <- test_matches %>%
  mutate(
    rf_pred_win = rf_prob > 0.5,
    rf_brier     = (1 - rf_prob)^2 # since actual label is 1 for "winner_name"
  )

misclass_rf <- mean(!test_matches$rf_pred_win[test_matches$season > 2014])
brier_rf    <- mean(test_matches$rf_brier[test_matches$season > 2014])

cat("RF misclassification rate:", round(misclass_rf,4), "\n")
```

```
## RF misclassification rate: 0.3512
```

```
cat("RF mean Brier score: ", round(brier_rf,4), "\n")
```

```
## RF mean Brier score: 0.2149
```

```

### CLUSTERING ###
library(factoextra) # for fviz_nbclust() & fviz_cluster()

## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa

train_test_matches <- matches %>% filter(season >= 2019)

# 1 & 2) stack winners and losers, flipping sign for losers
player_surface <- bind_rows(
  train_test_matches %>%
    select(surface, player = winner_name, all_of(covariates)),
  train_test_matches %>%
    select(surface, player = loser_name, all_of(covariates)) %>%
    mutate(across(all_of(covariates), ~ - .x))
)


# 3) average per player times surface
player_surface <- player_surface %>%
  group_by(player, surface) %>%
  summarise(
    across(all_of(covariates), ~ mean(.x, na.rm = TRUE)),
    n_matches = n(),
    .groups = "drop"
  ) %>%
  select(-n_matches)

# 4) pivot to wide: one row per player, cols = covariate_surface
player_wide <- player_surface %>%
  pivot_wider(
    id_cols = player,
    names_from = surface,
    values_from = all_of(covariates),
    names_sep = "_"
  ) %>%
  drop_na() # drop players missing any surface

# 5) scale features
feat_mat <- player_wide %>% select(-player)
feat_scaled <- scale(feat_mat)

# 6) elbow plot
fviz_nbclust(
  feat_scaled,
  kmeans,
  method = "wss",
  k.max = 25
) +
  labs(
    subtitle = "Elbow method",
    x = "Number of clusters k",
    y = "Total within-cluster sum of squares"
  )

```

```
# 7) run k-means
set.seed(123)
km <- kmeans(feats_scaled, centers = 10, nstart = 25)

# 8) attach cluster back
player_clusters <- player_wide %>%
  mutate(cluster = km$cluster)

# now `player_clusters` has columns:
#   player | diff_rolling_10_df_per_svgm_Clay | ... | diff_h2h_Hard | cluster
#
# can inspect each cluster's "specialty" by:
cluster_summary <- player_clusters %>%
  bind_cols(as_tibble(feats_scaled)) %>%
  group_by(cluster) %>%
  summarize(across(
    contains("svpt"),    mean, na.rm = TRUE
  ), across(
    contains("df_per_svgm"), mean, na.rm = TRUE
  ), across(
    contains("1stWon"),    mean, na.rm = TRUE
  ), across(
    contains("2ndWon"),    mean, na.rm = TRUE
  ))
```

```
## New names:
```

```

## * 'diff_rolling_10_df_per_svgm_Clay' -> 'diff_rolling_10_df_per_svgm_Clay...2'
## * 'diff_rolling_10_df_per_svgm_Grass' ->
##   'diff_rolling_10_df_per_svgm_Grass...3'
## * 'diff_rolling_10_df_per_svgm_Hard' -> 'diff_rolling_10_df_per_svgm_Hard...4'
## * 'diff_rolling_10_svpt_per_svgm_Clay' ->
##   'diff_rolling_10_svpt_per_svgm_Clay...5'
## * 'diff_rolling_10_svpt_per_svgm_Grass' ->
##   'diff_rolling_10_svpt_per_svgm_Grass...6'
## * 'diff_rolling_10_svpt_per_svgm_Hard' ->
##   'diff_rolling_10_svpt_per_svgm_Hard...7'
## * 'diff_rolling_10_1stWon_per_svgm_Clay' ->
##   'diff_rolling_10_1stWon_per_svgm_Clay...8'
## * 'diff_rolling_10_1stWon_per_svgm_Grass' ->
##   'diff_rolling_10_1stWon_per_svgm_Grass...9'
## * 'diff_rolling_10_1stWon_per_svgm_Hard' ->
##   'diff_rolling_10_1stWon_per_svgm_Hard...10'
## * 'diff_rolling_10_2ndWon_per_svgm_Clay' ->
##   'diff_rolling_10_2ndWon_per_svgm_Clay...11'
## * 'diff_rolling_10_2ndWon_per_svgm_Grass' ->
##   'diff_rolling_10_2ndWon_per_svgm_Grass...12'
## * 'diff_rolling_10_2ndWon_per_svgm_Hard' ->
##   'diff_rolling_10_2ndWon_per_svgm_Hard...13'
## * 'diff_rolling_10_bpSaved_per_bpFaced_per_svgm_Clay' ->
##   'diff_rolling_10_bpSaved_per_bpFaced_per_svgm_Clay...14'
## * 'diff_rolling_10_bpSaved_per_bpFaced_per_svgm_Grass' ->
##   'diff_rolling_10_bpSaved_per_bpFaced_per_svgm_Grass...15'
## * 'diff_rolling_10_bpSaved_per_bpFaced_per_svgm_Hard' ->
##   'diff_rolling_10_bpSaved_per_bpFaced_per_svgm_Hard...16'
## * 'diff_rolling_10_rank_Clay' -> 'diff_rolling_10_rank_Clay...17'
## * 'diff_rolling_10_rank_Grass' -> 'diff_rolling_10_rank_Grass...18'
## * 'diff_rolling_10_rank_Hard' -> 'diff_rolling_10_rank_Hard...19'
## * 'diff_rolling_10_ht_Clay' -> 'diff_rolling_10_ht_Clay...20'
## * 'diff_rolling_10_ht_Grass' -> 'diff_rolling_10_ht_Grass...21'
## * 'diff_rolling_10_ht_Hard' -> 'diff_rolling_10_ht_Hard...22'
## * 'diff_rolling_10_age_Clay' -> 'diff_rolling_10_age_Clay...23'
## * 'diff_rolling_10_age_Grass' -> 'diff_rolling_10_age_Grass...24'
## * 'diff_rolling_10_age_Hard' -> 'diff_rolling_10_age_Hard...25'
## * 'diff_h2h_Clay' -> 'diff_h2h_Clay...26'
## * 'diff_h2h_Grass' -> 'diff_h2h_Grass...27'
## * 'diff_h2h_Hard' -> 'diff_h2h_Hard...28'
## * 'diff_rolling_10_df_per_svgm_Clay' -> 'diff_rolling_10_df_per_svgm_Clay...30'
## * 'diff_rolling_10_df_per_svgm_Grass' ->
##   'diff_rolling_10_df_per_svgm_Grass...31'
## * 'diff_rolling_10_df_per_svgm_Hard' -> 'diff_rolling_10_df_per_svgm_Hard...32'
## * 'diff_rolling_10_svpt_per_svgm_Clay' ->
##   'diff_rolling_10_svpt_per_svgm_Clay...33'
## * 'diff_rolling_10_svpt_per_svgm_Grass' ->
##   'diff_rolling_10_svpt_per_svgm_Grass...34'
## * 'diff_rolling_10_svpt_per_svgm_Hard' ->
##   'diff_rolling_10_svpt_per_svgm_Hard...35'
## * 'diff_rolling_10_1stWon_per_svgm_Clay' ->
##   'diff_rolling_10_1stWon_per_svgm_Clay...36'
## * 'diff_rolling_10_1stWon_per_svgm_Grass' ->
##   'diff_rolling_10_1stWon_per_svgm_Grass...37'

```

```
## * 'diff_rolling_10_1stWon_per_svgm_Hard' ->
## 'diff_rolling_10_1stWon_per_svgm_Hard...38'
## * 'diff_rolling_10_2ndWon_per_svgm_Clay' ->
## 'diff_rolling_10_2ndWon_per_svgm_Clay...39'
## * 'diff_rolling_10_2ndWon_per_svgm_Grass' ->
## 'diff_rolling_10_2ndWon_per_svgm_Grass...40'
## * 'diff_rolling_10_2ndWon_per_svgm_Hard' ->
## 'diff_rolling_10_2ndWon_per_svgm_Hard...41'
## * 'diff_rolling_10_bpSaved_per_bpFaced_per_svgm_Clay' ->
## 'diff_rolling_10_bpSaved_per_bpFaced_per_svgm_Clay...42'
## * 'diff_rolling_10_bpSaved_per_bpFaced_per_svgm_Grass' ->
## 'diff_rolling_10_bpSaved_per_bpFaced_per_svgm_Grass...43'
## * 'diff_rolling_10_bpSaved_per_bpFaced_per_svgm_Hard' ->
## 'diff_rolling_10_bpSaved_per_bpFaced_per_svgm_Hard...44'
## * 'diff_rolling_10_rank_Clay' -> 'diff_rolling_10_rank_Clay...45'
## * 'diff_rolling_10_rank_Grass' -> 'diff_rolling_10_rank_Grass...46'
## * 'diff_rolling_10_rank_Hard' -> 'diff_rolling_10_rank_Hard...47'
## * 'diff_rolling_10_ht_Clay' -> 'diff_rolling_10_ht_Clay...48'
## * 'diff_rolling_10_ht_Grass' -> 'diff_rolling_10_ht_Grass...49'
## * 'diff_rolling_10_ht_Hard' -> 'diff_rolling_10_ht_Hard...50'
## * 'diff_rolling_10_age_Clay' -> 'diff_rolling_10_age_Clay...51'
## * 'diff_rolling_10_age_Grass' -> 'diff_rolling_10_age_Grass...52'
## * 'diff_rolling_10_age_Hard' -> 'diff_rolling_10_age_Hard...53'
## * 'diff_h2h_Clay' -> 'diff_h2h_Clay...54'
## * 'diff_h2h_Grass' -> 'diff_h2h_Grass...55'
## * 'diff_h2h_Hard' -> 'diff_h2h_Hard...56'

## Warning: There was 1 warning in 'summarize()'.
## i In argument: 'across(contains("svpt"), mean, na.rm = TRUE)'.
## i In group 1: 'cluster = 1'.
## Caused by warning:
## ! The '...' argument of 'across()' is deprecated as of dplyr 1.1.0.
## Supply arguments directly to '.fns' through an anonymous function instead.
##
## # Previously
## across(a:b, mean, na.rm = TRUE)
##
## # Now
## across(a:b, \(x) mean(x, na.rm = TRUE))
```

```
library(ggplot2)
library(ggrepel)
library(dplyr)
library(reshape2)

# PCA projection of your scaled features
pca_res <- prcomp(feats_scaled, center = TRUE, scale. = FALSE)

# build the scores data.frame
scores <- as.data.frame(pca_res$x[, 1:2])
colnames(scores) <- c("PC1", "PC2")
scores$player <- player_clusters$player
scores$cluster <- factor(player_clusters$cluster)
```

```

# generate distinct HCL hues
pal <- qualitative_hcl(10, palette = "Dark 3")

ggplot(scores, aes(PC1, PC2, color = cluster, label = player)) +
  geom_point(size = 2, alpha = 0.8) +
  geom_text_repel(
    max.overlaps = 25,
    size = 3,
    box.padding = 0.4,
    point.padding = 0.3
  ) +
  scale_color_manual(values = pal) +
  theme_minimal(base_size = 14) +
  labs(
    title = "Serve-Specialist Clusters on PC1 vs PC2",
    x = "PC1",
    y = "PC2",
    color = "Cluster"
  )

```

```

## Warning: ggrepel: 214 unlabeled data points (too many overlaps). Consider
## increasing max.overlaps

```

